Business case: The website went down at night



Nikita, DevOps engineer
Real case from my work. Client name hidden.

Business	An online service for checking receipts.
Risk	The site stopped at 02:41 at night. Nobody works at that time.
My work	Monitoring with alerts, a 20-minute check of the cause, a risk report.
Result	By 10:00 the team knew the cause. A hidden risk was found and a plan was made.

The business risk

- On 2 November 2026 at 02:41 at night, the website stopped working. Users could not open it.
- Before that, from about 23:30, the site had short drops of 1-2 minutes. They were too short for an alert.
- Without monitoring, a team often learns about such a failure from users in the morning.
- The team must know fast: is it our error, or a problem at the hosting provider?

What I did

- Monitoring checks the site all the time from a separate server.
- When the site was not available for more than 5 minutes, monitoring sent a critical alert to the team chat. Short drops did not create alerts, so the night chat was not full of noise.
- At 09:36 I started with the alert. I checked the domain from three public DNS services. The name servers of the provider itself also did not answer.
- I checked that the domain is paid. Then I checked the server address directly from five points in four countries.
- By 10:00 the picture was clear: the hosting provider had a full outage, both its DNS and its servers. Our setup had no error.
- I wrote a risk report: the domain, its settings and the server were all at one provider.

Site down at 02:41
users cannot open it

Alert at 02:46
after 5 minutes down

Checks 09:36-10:00
DNS, domain, server from 4 countries

Risk report
and a fix plan

Result

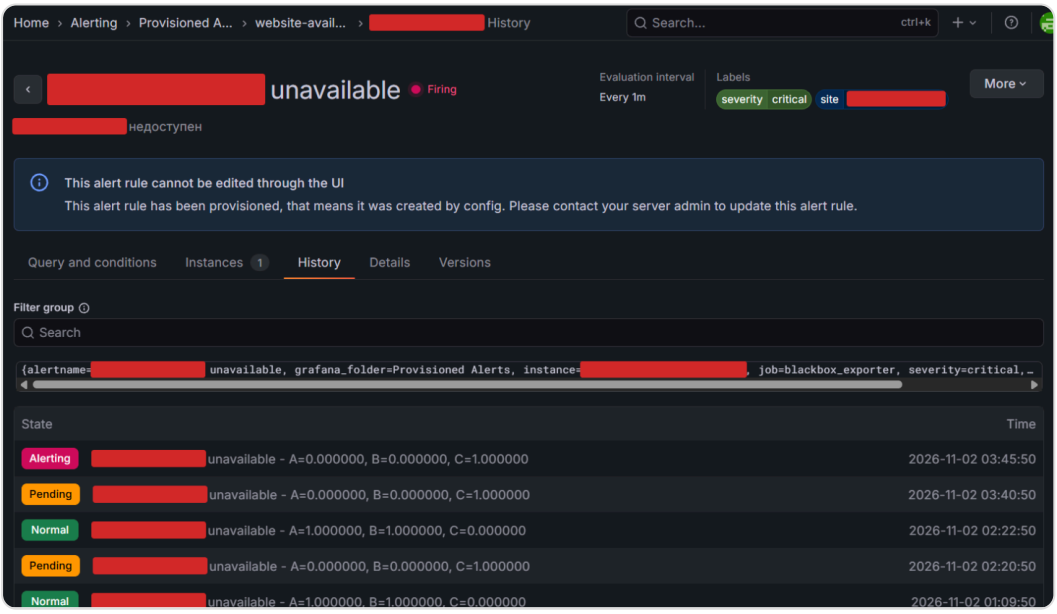
- The team knew about the failure from the alert, not from users.
- The cause was clear after about 20 minutes of checks, at 10:00.
- A hidden risk was found: one provider for everything. The plan: move domain settings and backups to other providers.

More results from my work

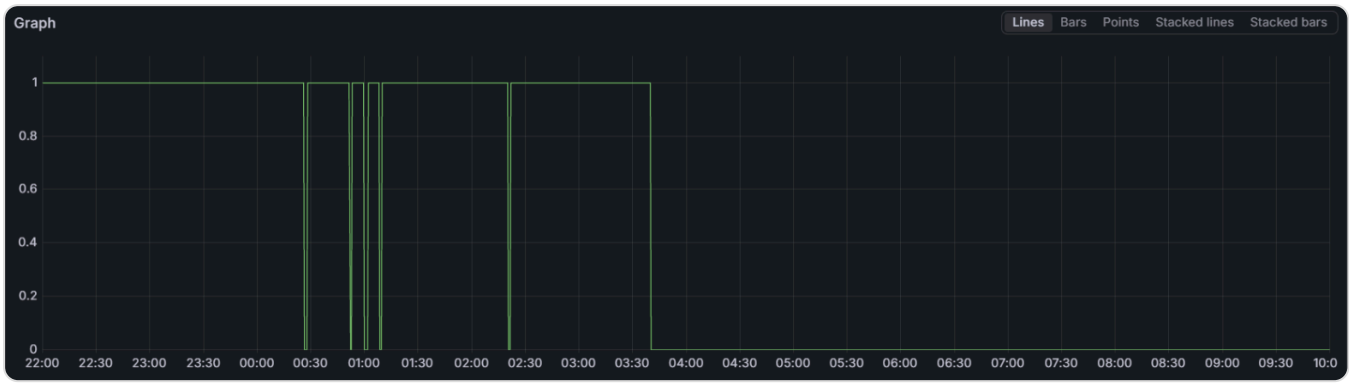
> **The GitLab server stopped 4 times in 2 months**
>
> It stopped on 4 August (twice), 9 September and 23 September and did not start again by itself. In the host logs I found the cause: the host gave its virtual machines almost twice more memory than it had (25.6 GB for VMs on a 15.5 GB host). On 23 September, 56 seconds before the stop, a new 4 GB VM was started. The same day I cut GitLab memory use from about 6 to about 3 GB, made other VMs smaller and added a 3.9 GB compressed swap. I also found password attacks on the control panel from the internet.

> **Alerts for 170+ websites**
>
> Each site has its own check. Failed backups and expiring certificates send an alert. I changed the limits and schedule. The team gets fewer false alerts.

How the result looks



The real alert history of that night. Real screenshot. Time is UTC+4: the alert fired at 02:46 Moscow time. The site name is covered in red.



Real availability graph of that night (UTC+4): 1 = site works, 0 = down. Short drops came first, then a full stop.



What you get: one screen with sites, certificate, processor, memory, disk. Real screenshot from my demo setup. Your names and numbers will differ.

What you get in this gig

- In this gig, I connect one Linux server to your existing Prometheus and Grafana.
- You get a dashboard with processor, memory and disk use.
- I set two alerts to your Telegram. For example: server not available and disk almost full.
- I test each alert in a safe way. You get short notes and steps to undo the changes.
- A new monitoring setup from zero is a separate offer after a check of your servers.

Words in this document

Monitoring	a system that watches your servers and sites all the time.
Alert	a message, for example in Telegram, when something goes wrong.
Dashboard	a screen with graphs that show the state of your servers.
Hosting provider	the company that rents you the server.

I do not show client names, server addresses or private data. I share them only with the client's permission. Icons: Tabler Icons, MIT license.